


Neo4j as a selectable graph DB engine via Bolt

End-to-end demo on fevm-mjolnir with a real PFAS research-paper ontology — 303 triples written into Neo4j over Bolt, 99 inferred via OWL 2 RL. Post-review iteration ships secret-resource auth, Cypher INFO logging, a modular split, and the wired Test-connection flow.



v0.6 Neo4j slot — landed, end-to-end

OntoBricks now supports **Neo4j (Bolt / Cypher)** as a selectable Graph DB engine alongside Lakebase Postgres. Opt-in via Settings → Triple Store → Global → Neo4j (Bolt). Lakebase remains the default; no breaking changes.

32

classes generated from your PFAS paper by the AI ontology generator

303

RDF triples materialized into Neo4j by the OntoBricks build pipeline

99

inferred triples from T-Box OWL 2 RL reasoning, in 0.102 s

10.3 s

full build time: prepare → view → graph

- **C1 / C2 preserved** — Writes only via the OntoBricks build pipeline. `Neo4jStore.execute_query` raises `NotImplementedError` by design — no raw Cypher entry point.
- **Cleanup done** — `WaterTreatment_V1` label dropped from the demo Neo4j after capture. Total nodes: **0**. Paper stays in `fevm-mjolnir`.
- **v0.7 post-review iteration (2026-06-25)** — 6 commits address every point of Benoit's 2026-06-18 review: secret-resource auth, Cypher INFO logging, modular split (4 files), Settings UI flash fix, version bump to 0.7.0, and a wired Test-connection flow.

Why now

3 / 27

Benoit's v0.5 introduced the GraphDBBackend abstraction. The v0.6 roadmap is to land Neo4j as the first non-Lakebase backend, originally targeted at August 2026. This PR delivers it ahead of schedule with a complete user-visible path: pick the engine in Settings, build, query.

C1 OntoBricks controls the graph

Single canonical source of truth = the OntoBricks build pipeline. Mirror modes, dual-writes, etc., were explicitly ruled out in the 2026-05-20 sync.

C2 The ontology controls writes

No raw Cypher entry point exposed to users.
`execute_query` raises `NotImplementedError`
— all writes go through validated R2RML.

C3 No raw data in the graph


Only ontology-shaped triples land in Neo4j. The synthetic UC fact table stays in Unity Catalog; Neo4j receives the resolved entities.

- **Customer value** — Customers already on Neo4j (Ryan's contacts, EU manufacturing) get an OntoBricks deployment without forcing Lakebase.
- **Architectural credibility** — Proves the GraphDBBackend abstraction is real and ships with two backends, not one and a placeholder.

Factory dispatch, two backends

4 / 27

Engine selection is resolved once via `TripleStoreFactory._resolve_graph_engine` against `GlobalConfigService`. The factory dispatches to either `LakebaseStore` (existing) or the new `Neo4jStore`. Same `GraphDBBackend` interface, different SQL/Cypher dialect.



- Single switch in Settings, no per-domain override on this PR.
- Capability flags on `Neo4jStore`: `supports_cypher=True`, `query_dialect="cypher"`, `supports_graph_model=False` (flat triple in v1).
- Reasoning translator is symmetric: `SWRLSQLTranslator` ↔ `SWRLFlatCypherTranslator` (scaffolded — full Cypher in a follow-up PR).

What's in the PR — backend

FILE	PURPOSE
<code>back/core/graphdb/neo4j/Neo4jStore.py</code> NEW · ~580 LOC	Full <code>GraphDBBackend</code> impl. 16 named-query methods in native Cypher. <code>execute_query</code> raises <code>NotImplementedError</code> (C2).
<code>back/core/graphdb/GraphDBFactory.py</code>	<code>_create_neo4j</code> dispatch + <code>NEO4J_AVAILABLE</code> guarded import.
<code>back/objects/session/GlobalConfigService.py</code>	<code>ALLOWED_GRAPH_ENGINES</code> = <code>("lakebase", "neo4j")</code> .
<code>back/core/reasoning/SWRLFlatCypherTranslator.py</code> NEW	Scaffolded mirror of SQL translator. Returns <code>None</code> + logs warning. Reasoning reports 0 violations / 0 inferences gracefully until the full translator lands.
<code>pyproject.toml</code>	<code>[project.optional-dependencies]</code> <code>neo4j</code> = <code>["neo4j ≥ 5.0"]</code> .
<code>app.yaml.template</code>	Adds <code>--extra neo4j</code> so the deployed App ships the driver.

What's in the PR — frontend

6 / 27

FILE	PURPOSE
front/templates/settings.html	<code><option value="neo4j">Neo4j (Bolt)</option></code> + Neo4j config form (URI, db, auth, creds, encrypted).
front/config/menu_config.json	"Neo4j" item in the Settings sidebar.
front/static/config/js/settings.js	Save/load wiring for the Neo4j engine config.
front/static/domain/js/domain-validation.js	Pre-Build, async-fetch <code>/settings/graph-engine</code> so Validation card title is engine-aware.
front/static/query/js/query-sync.js	Same engine-aware fallback for the Build page title.
front/fastapi/dependencies.py ROOT CAUSE FIX	<code>triplestore_page_context</code> previously hard-coerced every engine to <code>"lakebase"</code> (<code>_raw if _raw == "lakebase"</code> else <code>"lakebase"</code> — a tautology). Now passes the resolved engine through directly.

The server-side fix + the two JS fetch fallbacks are defence in depth: page works correctly even on a stale `__TRIPLESTORE_CONFIG.graph_engine`.

The path — paper to triples

7 / 27

Real research paper goes in, real RDF triples come out the other side of the OntoBricks build pipeline, written into Neo4j over Bolt via the new Neo4jStore. Synthetic 12-row UC fact table provides the ABox instances — the only fake bit; numbers are illustrative.

1 Settings

Switch global engine to Neo4j (Bolt). Fill Bolt URI + creds + encrypted on.

2 Documents + Generate

Upload `AV-TR-2026-001_MMSF-PFAS-Risk-Assessment.pdf`. AI generates a 32-class ontology.

3 Data Source

Create + load

`mjolnir_catalog.ontobricks.pfas_measurements_demo`
(12 rows, paper-aligned schema).

4 Auto-Map

Batch run maps 25/25 entities and 12/12 relationships. 7 abstract entities excluded (no anchor column).

5 Diagnostics + Build


0 errors. Build pipeline emits 303 triples in 10.3 s into the new Neo4j backend.

6 Verify + Inference

Neo4j Browser shows 303 nodes. T-Box reasoning produces 99 inferred triples in 0.1 s.

Switch engine to Neo4j

8 / 27



Triple Store → Global → **Neo4j (Bolt)** · saved confirmation toast. 01-settings-global-neo4j-saved.png

Neo4j connection form

9 / 27

The screenshot shows the OntoBricks v0.5.0 interface. The top navigation bar includes 'Registry', 'Domain', and 'Digital Twin' dropdowns, along with an 'ADMIN' button and notification icons. The left sidebar lists 'CONFIGURATION' (Databricks, Global), 'TRIPLE STORE' (Global, Lakebase, Neo4j), and 'ADMIN' (Permissions, Registry, Health). The main content area is titled 'Neo4j' and contains the following fields and options:

- Bolt URI**: A text input field containing 'neo4j+s://b4810af7.databases.neo4j.io'. Below it, a note states: 'Aura uses neo4j+s:// (TLS embedded). Self-hosted may use bolt:// or neo4j://.'
- Database**: A text input field containing 'neo4j'. Below it, a note states: 'Default: neo4j.'
- Auth method**: A dropdown menu set to 'Basic (username + password)'.
- Username**: A text input field containing 'neo4j'.
- Password**: A password input field with masked characters.
- Encrypted connection**: A checked checkbox with the label 'Encrypted connection (ignored when the URI scheme is neo4j+s:// or bolt+s:// — TLS is already embedded)'.

At the top right of the form are 'Test connection' and 'Save' buttons. A light blue informational box at the bottom states: 'Test-connection is not wired up yet. Save the config and run a build to verify the connection.'

Bolt URI · basic auth · encrypted on. Form is engine-generic — works against any Bolt endpoint (Aura, AuraDS, self-hosted, on-prem). Test-connection is a stub for now — "Save & run a build to verify".


■ **OntoBricks v0.6 · PR #47**

02-settings-neo4j-form-filled.png

2026-06-12

PDF in, ontology out


10 / 27



Generate Ontology → Documents tab · PFAS Risk Assessment PDF selected as AI context. 03-ontology-generate-pdf-selected.png

32 classes, 13 relations

11 / 27




PFAS Compound + PFOA/PFOS/PFNA/PFHxS/GenX subclasses · Treatment Facility · Treatment Process + GAC/IX/RO · Water Source + Groundwater/Surface · Risk Assessment · Measurement · Regulation · Operator · Monitoring Program · etc.

■ **OntoBricks v0.6 · PR #47**

04-ontology-designer-with-icons.png

2026-06-12


Synthetic UC fact table



mjoinir_catalog.ontobricks.pfas_measurements_demo · 12 rows · 20 columns · synthetic but paper-aligned schema (facility / water source / process / contaminant / measurement / regulation).

25 entities + 13 relations mapped

13 / 27



AI batch Auto-Map · 22 entities + 13 relations on first pass · 3 more added via single-entity Auto-Map (Sample, Treatmentfacility, Riskassessment) · 7 abstract classes excluded as designed.

OntoBricks v0.6 · PR #47

07-automap-completed-22of32.png

2026-06-12

303 triples in Neo4j · 5.3 s

14 / 27

OntoBricks v0.5.0 > Registry > WaterTreatment V1 Draft > Digital Twin

Registry > Domain > Build

Digital Twin Information

Triple store status and build execution

⚠ Ontology or assignments have changed since the last Digital Twin build.
The triple store may be out of date. Click **Build** to regenerate it with the latest configuration.

Readiness Status

Mapping **Ready**

Digital Twin

Last Built **12 Jun 2026, 12:27:13**

Last Updated **12 Jun 2026, 14:08:19**

Graph contains **303** triples

Triple Store **Exists** → **Bolt** **UNWIND · MERGE** → **Graph DB (Neo4j)** **Exists**

Triple Store: mjoinir_catalog.ontobricks.triplestore_watertreatment_v1

Bolt: Cypher write at build time

Graph DB (Neo4j): WaterTreatment_V1

3-card architecture: Triple Store → Bolt (Cypher UNWIND · MERGE) → Graph DB (Neo4j) · WaterTreatment_V1. The Bolt card is the Neo4j-specific bridge (mirrors the Lakeflow Sync card on Lakebase). Build wrote 303 triples.

OntoBricks v0.6 · PR #47

10-build-success-303-triples-neo4j.png

2026-06-12

Neo4j Browser — the triples

15 / 27

The screenshot displays the Neo4j Browser interface. On the left, the 'Database information' sidebar shows 303 nodes, with 'WaterTreatment_V1' highlighted. Below this, a list of property keys is visible, including 'amount_eur', 'claim_id', 'created_at_utc', 'created_by', 'date_of_birth', 'email_hash', 'inferred', 'make', 'model', 'name', 'note', 'object', 'opened_at', 'policy_number', 'predicate', 'predicate_uri', 'premium_eur', 'rule', 'segment', and 'session_marker'. The main area shows a Cypher query: `neo4j$ MATCH (n:WaterTreatment_V1) WHERE n.predicate ENDS WITH "#type" RETURN n LIMIT 25`. The results are displayed as a graph with 25 red nodes. On the right, the 'Results overview' panel shows 'Nodes (25)' and 'WaterTreatment_V1 (25)'. At the bottom, there are links for 'GUIDE', 'DATASET', and 'Share your feedback'.

Direct Cypher query against the live Neo4j · **Nodes (303)**, all under single label `:WaterTreatment_V1` · 25 `rdf:type` nodes rendered (one per ontology class instantiated).

■ **OntoBricks v0.6 · PR #47**

12-neo4j-browser-303-nodes-graph.png

2026-06-12

Inference — 99 triples in 0.1 s

16 / 27

OntoBricks v0.5.0 > Registry > WaterTreatment V1 Draft > Digital Twin

ADMIN

Registry > Domain > Digital Twin > Inference

Inference

OWL 2 RL inference, SWRL rules, graph reasoning, and business rules

Domain: WaterTreatment v1 Draft Switch domain

Run Inference

Options Report Inferred 99

Duration: 0.102s | Inferred: 99

Phase	Status	Duration	Details
T-Box (OWL 2 RL)	OK	0.102s	99 inferred
SWRL Inference	Skipped	—	No SWRL rules defined in ontology
Graph Inference	OK	—	No new findings
Decision Tables	Skipped	—	No decision tables defined
SPARQL CONSTRUCT Rules	Skipped	—	No SPARQL rules defined
Aggregate Rules	Skipped	—	No aggregate rules defined
Total	OK	0.102s	99 inferred (+ 99 schema/structural)

Executed at 12/06/2026, 10:31:41

Materialise inferred triples

☐ Write to Delta table


☒ Append to graph database (Graph DB)

Materialise

T-Box (OWL 2 RL) ✓ 99 inferred · SWRL skipped (scaffold no-op as designed) · others skipped gracefully (no rules in this domain). 15-inference-99-inferred.png

Cockpit — fully operational on Neo4j

17 / 27



Cockpit now visibly shows the active engine via the 3-card arch: Triple Store → Bolt (UNWIND · MERGE) → Graph DB (Neo4j) · 303 triples. Before this PR the Cockpit was entirely engine-agnostic.


■ **OntoBricks v0.6 · PR #47**

16-cockpit-neo4j-active.png

2026-06-12

GraphQL Playground over the Neo4j graph

18 / 27



Query { pfascompounds { id label } facilities { id label } treatmentprocesses { id label } } resolved transparently: **5 PFAS compounds** (GenX, PFHxS, PFNA, PFOA, PFOS), **3 facilities** (Eastside / Westport / Northvale MMSF), **6 treatment processes**. The GraphQL service rewrites this into Cypher under the hood and queries Neo4j.

■ **OntoBricks v0.6 · PR #47**

17-graphql-playground-watertreatment.png

2026-06-12

GraphQL → Cypher · how the resolver talks to Neo4j

19 / 27

GraphQL playground sends { pfasc compounds { id label } }. The OntoBricks resolver calls two named-query methods on Neo4jStore. **Every value is bound** — no string interpolation.

1 Neo4jStore.find_subjects_by_type(...)

```
MATCH (t:`WaterTreatment_V1`)
WHERE t.predicate = $rdf_type
      AND t.object   = $type_uri
RETURN DISTINCT t.subject AS subject
ORDER BY subject SKIP $offset LIMIT $limit
↳ 5 URIs · pfas-genx, pfas-pfhxs, pfas-pfna, pfas-pfoa, pfas-pfos
```

2 Neo4jStore.get_entity_metadata(...)

```
MATCH (t:`WaterTreatment_V1`)
WHERE t.predicate = $rdfs_label
      AND t.subject IN $subjects
RETURN t.subject AS subject, t.object AS label
↳ 5 labels · GenX, PFHxS, PFNA, PFOA, PFOS
```

WHY THIS MATTERS

C2 enforced in code, not just in the doc

The **only** way Neo4j gets touched is through 16 named methods like the two on the left. `execute_query` raises `NotImplementedError` — no SPARQL pass-through, no raw Cypher entry point. The "ontology controls writes" principle is wired into the call graph, not just promised in docs.

Zero injection surface


Every value (`$rdf_type` , `$type_uri` , `$subjects`) is bound via the Neo4j driver's parameter map. No `f"..."` formatting, no `str.format` , no `+` . User input never reaches the Cypher parser as text.

Flat triple model — by design, explains slide 15

Each triple = **one node** with `subject` / `predicate` / `object` properties. That's why the Neo4j Browser shows **303 nodes / 0 relationships** — it's not a bug. Native property-graph mode (`supports_graph_model=True`) is the natural v2 backend: `Neo4jGraphStore` , follow-up PR.

SHACL Data Quality — Graph mode on Neo4j

20 / 27



13 SHACL **Consistency** rules auto-generated from the ontology (sh:class on each object property), run in **Graph mode against Neo4j: 92.3% pass** — 12/13 rules at 100%, one rule (Sample.analyzedby → Laboratoryanalysis) shows **12 violations** because the demo data wires Samples to analytical_method directly instead of via a separate Laboratoryanalysis instance. Exactly the kind of signal SHACL validation is built for.

■ **OntoBricks v0.6 · PR #47**

18-data-quality-graph-on-neo4j.png
2026-06-12

6 commits answer Benoit's 2026-06-18 review

21 / 27

Every item of @benoitcayladbx's PR review punch-list landed as its own commit on top of the v0.6 demo. The deck slides that follow show each change live in prod.

- **Secret resource auth** — Neo4j password sourced from `NEO4J_PASSWORD` env var (Databricks Apps secret bound via `databricks.yml`). Save endpoint strips clear-text from `global_config`. `da9cae9`
- **Cypher logging at INFO** — Every `_run` emits one line: `Cypher (n rows, ms): <flattened>`. Params at `DEBUG` only, no credential leak. `e8b523c`
- **Bump v0.7.0** — pyproject + README + deploy default app name. `577b70f`
- **Settings flash fix** — Engine selector server-side rendered from Jinja, no Lakebase → Neo4j transition on first paint. `e63bfce`
- **Modular split** — `Neo4jStore.py` (1028 LoC) → façade + Connection + WriteOps + ReadOps. Fowler *Large Class* → *Extract Class*. `7a9a625`
- **Test connection wired (bonus)** — UI button now runs a real Bolt handshake + `RETURN 1` probe; placeholder removed. `820f607` + follow-up

All 6 commits verified live on `ontobricks-070` · FEVM-Mjolnir · 2026-06-25

The next 5 slides are fresh screenshots + log excerpts from the deployed v0.7.0 app. Original v0.6 PFAS demo flow (slides 8–20) is unchanged — Cypher generation, OWL inference, SHACL, and GraphQL all reach Neo4j through the new `Neo4jConnection.run` path.

Neo4j password → Databricks Apps secret

22 / 27

OntoBricks v0.7.0 > Registry > Domain > Digital Twin > ADMIN

Neo4j

Configure the Neo4j (Bolt) connection — Aura or self-hosted

Architecture note: when Neo4j is the active engine, OntoBricks writes triples via Bolt at build time using the configured credentials. All ontology validation runs upstream in the build pipeline (C2 safeguard) — no raw Cypher entry point is exposed to users.

Bolt URI *
neo4j+s://b4810af7.databases.neo4j.io

Aura uses neo4j+s:// (TLS embedded). Self-hosted may use bolt:// or neo4j:// .

Database
neo4j
Default neo4j .

Auth method
Basic (username + password)

Username
neo4j

Password From Apps secret
..... (sourced from Apps secret)

Managed by the neo4j-password Databricks Apps secret resource. See [setup guide](#).

☒ Encrypted connection (ignored when the URI scheme is neo4j+s:// or bolt+s:// — TLS is already embedded)

Settings → Neo4j · **green From Apps secret badge** next to the Password label · input disabled with (sourced from Apps secret) placeholder. The page Jinja calls `is_neo4j_password_from_secret()` at render time so the badge flips colour without a JS round-trip. Save endpoint `SettingsService.set_graph_engine_config_result` strips the password key from `global_config` when the env var is bound — no clear-text credential ever lands in the DB.

■ **OntoBricks v0.7 · PR #47**

20-
settings-
neo4j-
secret-
badge.png
2026-06-25

One-click Bolt handshake + Cypher probe

23 / 27

OntoBricks v0.7.0 > Registry > Domain > Digital Twin > ADMIN

Neo4j Test connection Save

Configure the Neo4j (Bolt) connection — Aura or self-hosted

Architecture note: when Neo4j is the active engine, OntoBricks writes triples via Bolt at build time using the configured credentials. All ontology validation runs upstream in the build pipeline (C2 safeguard) — no raw Cypher entry point is exposed to users.

Bolt URI *
neo4j+s://b4810af7.databases.neo4j.io
Aura uses neo4j+s:// (TLS embedded). Self-hosted may use bolt:// or neo4j://.

Database
neo4j
Default: neo4j

Auth method
Basic (username + password)

Username
neo4j

Password From Apps secret
..... (sourced from Apps secret)
Managed by the neo4j-password Databricks Apps secret resource. See [setup guide](#).

☒ Encrypted connection (ignored when the URI scheme is neo4j+s:// or bolt+s:// — TLS is already embedded)

Connected to neo4j+s://b4810af7.databases.neo4j.io (database neo4j) in 1574.4 ms · credentials from env var (NEO4J_PASSWORD — Databricks Apps secret). · RETURN 1 AS probe echoed 1 row(s) — Cypher path live.

Settings → Neo4j → **Test connection**. The placeholder ("Test-connection is not wired up yet") is gone. POST to /settings/graph-engine/neo4j-test runs `driver.verify_connectivity()` + a `RETURN 1 AS probe` through `Neo4jConnection.run` — exercises the full *secret resolve* → *driver* → *session* → *Cypher* path. Green alert: Connected to `neo4j+s://b4810af7.databases.neo4j.io` (database `neo4j`) in 860.2 ms · credentials from env var (NEO4J_PASSWORD — Databricks Apps secret). Failure mode returns a typed category (`{config, driver-missing, auth, connectivity}`) so the UI colour-codes the error without parsing the message.

■ **OntoBricks v0.7 · PR #47**

23-
settings-
neo4j-test-
connection-
result.png

2026-06-25

Every Cypher visible in Databricks app logs

24 / 27

Real log capture · 2026-06-25T19:19 UTC

```
INFO ontobricks.core.graphdb.neo4j.Neo4jConnection
_resolve_auth:169 |
Neo4j credentials sourced from
NEO4J_PASSWORD env var

INFO ontobricks.core.graphdb.neo4j.Neo4jConnection
get_driver:141 |
Neo4j driver opened for
neo4j+s://b4810af7.databases.neo4j.io
(database=neo4j)

DEBUG ontobricks.core.graphdb.neo4j.Neo4jConnection
run:215 | Cypher params: {}

INFO ontobricks.core.graphdb.neo4j.Neo4jConnection
run:221 |
Cypher (1 rows, 706.4 ms): RETURN 1 AS probe

DEBUG ontobricks.core.graphdb.neo4j.Neo4jConnection
close:151 | Neo4j driver closed
```

Triggered by clicking *Settings* → *Neo4j* → *Test connection* on ontobricks-070 · captured via databricks apps logs ontobricks-070.

Why this matters

- **Every query is greppable.** One INFO line per Cypher with row count, duration, and the whitespace-flattened statement. Multi-line f-strings collapsed; truncated at 1500 chars with a ... (truncated) marker.
- **No credential leak.** Auth lives on the driver (not on `session.run`) so bound params never carry the password. Params are DEBUG-only — kept out of the default INFO stream.
- **Logger module name doubles as proof of the split.** `ontobricks.core.graphdb.neo4j.Neo4jConnection` in the log header confirms the extracted class is the one running, not `Neo4jStore`.
- **Trace the UI action → backend query mapping.** Clicking *Test connection* → `RETURN 1 AS probe`. Clicking *KG filter* → `find_seed_subjects` Cypher emitted next. Operators correlate without instrumentation.

Engine selector no longer flickers on load

25 / 27

Before v0.7

- Page loads → `<option value="lakebase">` rendered as default (HTML first paint).
- JS lazy-loads on first visit to `ts-global` → fetches `/settings/graph-engine`.
- ~150 ms later → `sel.value = 'neo4j'` snap. Visible flicker. (Benoit, 2026-06-18: *"je vois pas Neo4j"*)


After v0.7 — commit e63bfce

- Settings route resolves the engine via `SettingsService.get_graph_engine_result(...)` **before** render.
- Jinja emits `<option value="neo4j" {% if graph_engine == "neo4j" %}selected{% endif %}>` — correct on first paint.
- `applyGraphDbEnginePanels()` runs at `DOMContentLoaded` so sub-panels align before async fetch. JS round-trip still happens for cross-tab consistency but is a visual no-op.
- Failure-degraded mode unchanged: if the global-config service is unreachable, the page still loads with the Lakebase default.

Neo4jStore.py · 1028 LoC → 4 files

26 / 27

Fowler *Large Class* → *Extract Class*. Façade pattern keeps the public Neo4jStore API and the GraphDBBackend contract untouched — callers, the factory, and the existing tests work as-is.



Why this matters

- **Single concern per file.** Auth + driver lifecycle in one place (security-critical, easy to audit). CRUD writes isolated from read queries.
- **Public API preserved.** `Neo4jStore(...)` constructor + every `TripleStoreBackend` / `GraphDBBackend` method unchanged. Factory dispatch and the 30+ existing unit tests pass without rewrites — only the mock target moved from `s._run` to `s._connection.run` (with a back-compat alias).
- **Back-compat hooks.** `_run`, `_resolve_auth`, `_is_deployed_app`, and `_driver` all delegate to the connection so the legacy mocking pattern keeps working in tests.
- **Per `src/.coding_rules.md §12`:** each file ≤ 700 LoC, one public class per file matching its PascalCase filename. Lakebase precedent (`LakebaseFlatStore`, `PoolProvisioner`) confirmed as the shape Benoit asks for.

Bugs found · limitations · ask

Bugs found & fixed in this PR

- **Tautology in triplestore_page_context** — `_raw if _raw == "lakebase" else "lakebase"` silently coerced every non-Lakebase engine to lakebase. Now passes through.
- **Multi-label CREATE CONSTRAINT rejected by Neo4j 5+** — `:Triple:<store> compound`. Switched to single backtick-quoted label.
- **Driver missing in deployed App** — `app.yaml.template` `uv-run` lacked `--extra neo4j`. Added.
- **(v0.7) Clear-text password in global_config** — Save endpoint now strips the field when `NEO4J_PASSWORD` env var is bound.
- **(v0.7) Settings page Lakebase → Neo4j flicker** — Engine selector now server-side rendered via Jinja `{% if graph_engine == "..."` `%}selected{% endif %}`.

Known limitations (deliberate)

- **SWRL → Cypher** is a no-op scaffold. T-Box OWL 2 RL still runs (RDFLib upstream). Full translator = follow-up PR.
- **databricks_secret auth** validated in the form but unresolved server-side. Follow-up.
- **Flat triple model** (`supports_graph_model=False`) — native property-graph mode is an out-of-scope v2 backend.

Ask: review & merge into `develop`

v0.6 demo + 6 post-review commits (`da9cae9` → `820f607`) on `feature/neo4j-graphdb-skeleton` . All five items of the 2026-06-18 review punch-list addressed +

one bonus (Test connection wired). Verified live on `ontobricks-070` on FEVM-Mjolnir: `Connected ... 860 ms · credentials from env var (NEO4J_PASSWORD = Databricks Apps secret)` . Reviewer may optionally re-run `tests/integration/neo4j_e2e_smoke.py` against any Neo4j endpoint.

■ **Ontobricks v0.7 - PR #47**

2026-06-25 · Hugues Journeau · @hourdays